



Kubernetes From Scratch – Complete Guide

Welcome to this comprehensive, visually enriched guide to Kubernetes (K8s)!
This guide uses help you understand container orchestration step-by-step.

Let's orchestrate some magic! 🚀



What is Kubernetes?

Kubernetes is an **open-source container orchestration platform** that automates the deployment, scaling, and management of containerized applications across clusters. Originally developed by Google, it is now maintained by the **Cloud Native Computing Foundation (CNCF)**.



Key Benefits

<u>Feature</u>	<u>Description</u>
High Availability	Ensures applications run redundantly across nodes to reduce downtime.
Zero-Downtime Deployments	Performs rolling updates without service interruption.
Auto-Scaling	Adjusts resources automatically based on metrics (CPU, memory).
Auto-Healing	Restarts failed containers, reschedules pods, and maintains replicas.



What Does Container Orchestration Do?

Kubernetes handles the heavy lifting of managing containers at scale.

Core Functions

- **Provision & Deploy** containers across multiple nodes
- **High Availability** through balanced distribution
- **Auto-Scaling** workloads (via HPA, VPA)
- **Efficient Resource Allocation** (CPU, memory, storage)
- **Load Balancing & Traffic Routing**
- **Service Discovery** between microservices
- **Secure Communication** via Network Policies


Popular Orchestrators Comparison

Orchestrator	Strengths	Best For	Drawbacks
Kubernetes	Feature-rich, flexible, huge ecosystem	Enterprise apps	Steep learning curve
Docker Swarm	Simple, Docker-native	Small teams, quick deployments	Less advanced
Apache Mesos	Handles mixed workloads	Big data clusters	Complex, declining adoption

Kubernetes Components & Architecture

Kubernetes operates through a cluster — a group of machines working together.

Cluster Overview



Master Node = Control plane = Brain

Worker Nodes = Muscle

Multi-Master Setup = High Availability

■ Master Node Components

Component	Role	Key Features
API Server	Central endpoint for all REST/K8s requests	Handles kubectl, validates config, exposes /metrics
Controller Manager	Maintains desired state	Includes node/replica controllers
Scheduler	Assigns pods to nodes	Considers resources, affinities, taints
etcd	Distributed key-value store	Stores entire cluster state

Essential Add-ons

- **CoreDNS** – Service discovery
- **Kubernetes Dashboard** – Web UI
- **Container Runtime** – Docker, containerd, CRI-O

■ Worker Node Components

<u>Component</u>	<u>Role</u>	<u>Key Features</u>
Kubelet	Node agent managing pod lifecycle	Runs containers, performs health checks
Kube-Proxy	Network routing for services	Implements load balancing via iptables/IPVS

Pro Tip: Kubelet communicates only with the API Server, never directly with other nodes.

✂ kubectl — The Swiss Army Knife CLI

General syntax:

```
kubectl [command] [type] [name] [flags]
```

Examples of Commands

- **Commands:** get, apply, create, delete, describe
- **Types:** pod, deployment, service, namespace
- **Flags:** -n, -o yaml, --all-namespaces

kubectl Cheat Sheet

Action	Command Example
List resources	<code>kubectl get pods -A</code>
Describe	<code>kubectl describe pod my-pod</code>
Create/Apply	<code>kubectl apply -f deployment.yaml</code>
Scale	<code>kubectl scale deploy/my-app --replicas=5</code>
Logs & Exec	<code>kubectl logs my-pod</code> <code>kubectl exec -it my-pod -- /bin/sh</code>
Delete	<code>kubectl delete pod my-pod</code>
Port-forward	<code>kubectl port-forward svc/my-svc 8080:80</code>
Events	<code>kubectl get events --sort-by=.metadata.creationTimestamp</code>

Core Kubernetes Concepts

Pod — The Smallest Unit

Pods wrap around one or more containers that share:

- Network namespace (one IP)
- Storage volumes
- Lifecycle management rules

Pod Types

- **Single-container Pod**
 - **Multi-container Pod** (Sidecar pattern)
-

Deployment — Declarative Scaling

- Manages ReplicaSets
- Handles rolling updates/rollbacks
- Ensures stateless workloads scale

Deployment YAML

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: nginx
          image: nginx:1.14
```

Service — Stable Networking

<u>Type</u>	<u>Use Case</u>	<u>Exposure</u>
ClusterIP	Internal-only	Virtual IP inside cluster
NodePort	External access via nodes	<NodeIP>:<NodePort>
LoadBalancer	Cloud LB	Public external IP

Namespace — Virtual Clusters

Used for isolation, organization, and resource limits.

Default namespaces:

- default
 - kube-system
 - kube-public
-

Secrets, ConfigMaps & Storage

<u>Concept</u>	<u>Purpose</u>	<u>Security Notes</u>
Secrets	Store sensitive info	Base64 encoded, not encrypted
ConfigMaps	Store config data	Mounted as env or volumes
Volumes	Persist data	Supports emptyDir, hostPath, PVC

ReplicaSet & DaemonSet

- **ReplicaSet** – Maintains N identical Pods
 - **DemonSet** – Runs one Pod per node (logs, monitoring)
-

Labels, Selectors & Annotations

- **Labels:** metadata to organize resources
 - **Selectors:** query/filter resources
 - **Annotations:** metadata for external tools
-

Affinity, Anti-Affinity & Taints

Control where Pods are scheduled:

- **Node Affinity** – Match node labels
- **Pod Affinity/Anti-Affinity** – Group or separate pods
- **Taints/Tolerations** – Restrict scheduling on specific nodes

RBAC — Role-Based Access Control

<u>Resource</u>	<u>Scope</u>	<u>Purpose</u>
Role	Namespace	Defines permissions
ClusterRole	Cluster	Global permissions
RoleBinding	Namespace	Binds Role
ClusterRoleBinding	Cluster	Binds ClusterRole

Best Practice: Always follow **least privilege**.

Networking & Ingress

Ingress — Smart HTTP Routing

Feature	Benefit
Path-based routing	/api → backend1, /web → frontend
TLS termination	HTTPS security
Host-based routing	api.example.com, app.example.com

Scaling & Autoscaling

Horizontal Pod Autoscaler (HPA)

Automatically scales deployments.

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: my-app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-app
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```

Vertical Pod Autoscaler (VPA)

Adjusts CPU/memory requests automatically (may restart pods).



Monitoring, Logging & Troubleshooting

Built-in Tools

- `kubectl get events -w`
- `kubectl logs my-pod -f`
- `kubectl top pods`

Observability Stack

<u><i>Tool</i></u>	<u><i>Purpose</i></u>
Prometheus	Metrics
Grafana	Dashboards
EFK Stack	Centralized logs
Jaeger	Tracing

Troubleshooting Flow:

- `kubectl describe pod <name>`
 - `kubectl exec -it <pod> -- curl <service>`
-



Security Best Practices

- Enforce Pod Security Standards
- Default-deny Network Policies
- Image scanning (Trivy, Clair)
- External Secret Stores
- Admission Controllers
- API server audit logs